

Assignment Project Exam Help

Recovery

<https://powcoder.com>

P.J. McBrien

Imperial College London


Add WeChat powcoder

DBMS Architecture

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Recovery Manager (RM)

Assignment Project Exam Help

protect the DBMS against failures

- **system failures** loss of volatile storage

- 1 committed transactions written to disc
- 2 uncommitted transactions not written to disc

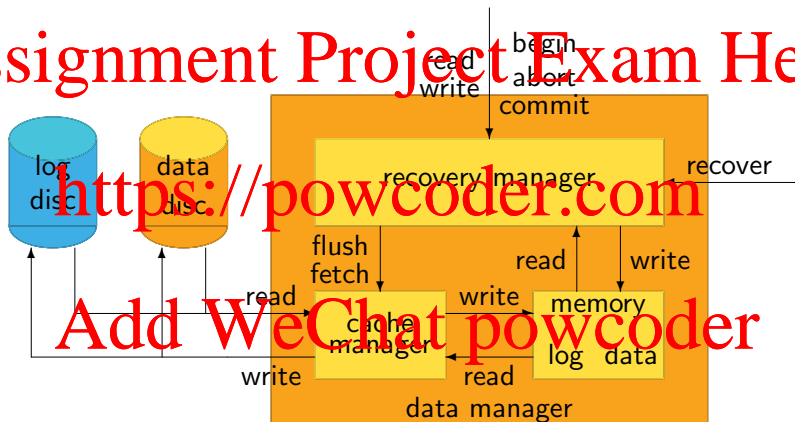
OR

- 3 sufficient information such that (1) and (2) may be met by a *recovery*

- **media failures** loss of stable storage

Add WeChat powcoder

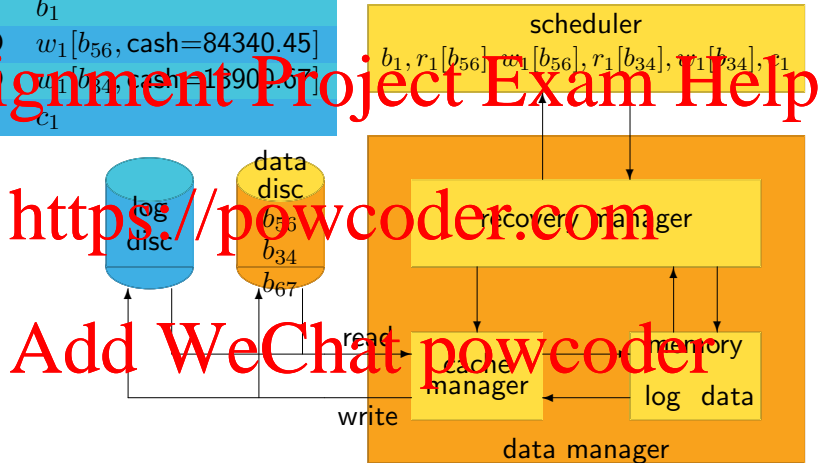
Enhanced Data Manager Architecture



- Need to cache log as well

Need to REDO


| | |
|------|-------------------------------------|
| LOG | b_1 |
| REDO | $w_1[b_{56}, \text{cash}=84340.45]$ |
| REDO | $w_1[b_{34}, \text{cash}=13900.57]$ |
| LOG | c_1 |



- REDO required if committed transactions not in stable storage
- must write all REDO to log before commit of transaction

Need to UNDO

LOG b_1
 UNDO $w_1[b_{56}, \text{cash}=94340.45]$
 UNDO $w_1[b_3, \text{cash}=8900.67]$
 LOG c_1



- UNDO required if non-committed transactions in stable storage
- Must flush UNDO to log before corresponding write to data

Quiz 1: Contents of Data Disc After a Transaction

| branch | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

```

BEGIN TRANSACTION T1
UPDATE branch
SET cash=cash-10000.00
WHERE sortcode=56

UPDATE branch
SET cash=cash+10000.00
WHERE sortcode=34
COMMIT TRANSACTION T1
  
```

| branch ① | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ③ | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 84340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ② | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ④ | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 84340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

What must the contents of the branch table on the data disc be after the transaction commits?

A

④

B

① or ④

C

①, ③ or ④

D

①, ②, ③ or ④

Quiz 2: Contents of Log Disc After a Transaction

Data Disc Before Transaction

| branch | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

BEGIN TRANSACTION T1

```
UPDATE branch
SET cash=cash-10000.00
WHERE sortcode=56
```

```
UPDATE branch
SET cash=cash+10000.00
WHERE sortcode=34
```

COMMIT TRANSACTION T1

Data Disc At Commit Time

| branch | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

<https://powcoder.com>

What must be on the log disc after commit time?

A

REDO r_{56} REDO r_{34} UNDO r_{56} UNDO r_{34}

B

REDO r_{56} UNDO r_{34}

C

UNDO r_{34}

D

REDO r_{56}

Before and after images

before image

| | branch | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

$w_1[b_{56}]$

Add WeChat powcoder

| | branch | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 84340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

after image

- before image allows RM to **undo** $w_1[b_{56}]$
- after image allows RM to **redo** $w_1[b_{56}]$

Database Logs

Assignment Project Exam Help

```
LOG      b1
REDO    w1[b56, cash=84340.45]
REDO    w1[b34, cash=18900.67]
LOG      c1
```

```
LOG      b1
UNDO    w1[b56, cash=94340.45]
UNDO    w1[b34, cash=8900.67]
LOG      c1
```

<https://powcoder.com>

Add WeChat powcoder

```
LOG      b1
UNDO    w1[b56, cash=94340.45]
REDO    w1[b56, cash=84340.45]
UNDO    w1[b34, cash=8900.67]
REDO    w1[b34, cash=18900.67]
LOG      c1
```

What must a complete REDO/UNDO log contain?

Must contain

- REDO information for each update
- UNDO information for each update
- commit of each transaction

Might contain

- begin of each transaction
 - can be inferred from first REDO/UNDO
 - presence useful to stop search of UNDO records
- abort of each transaction
 - can be inferred from lack of commit
 - presence useful to indicate UNDO already done

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Rules for log and data updates

Assignment Project Exam Help

write ahead logging (WAL)

Redo rule

- commit \rightarrow flush log of transaction to disc
- never respond to scheduler before log written

Undo rule:

- flushing uncommitted data \rightarrow flush log of operations

<https://powcoder.com>

Add WeChat powcoder

Basic Recovery Procedure

Assignment Project Exam Help



$\Rightarrow w_v[o_1], c_v, w_z[o_2], u_y[o_1], c_y, w_z[o_2] \Rightarrow$



- 1 UNDO — *Scan back through the log*
 - Collect set of committed transactions $C = \{v, y\}$
 - Collect set of incomplete transactions $I = \{x, z\}$
 - Perform UNDO for any transaction in $I = w_z[o_2], w_x[o_2]$
- 2 REDO — *Scan forward through the log*
 - Perform REDO for any transaction in $C = w_v[o_1], w_y[o_1]$



$\Rightarrow \text{UNDO}(w_z[o_2]), \text{UNDO}(w_x[o_2]), \text{REDO}(w_v[o_1]), \text{REDO}(w_y[o_1]) \Rightarrow$



Example of Recovery

Log

| | |
|------|-------------------------------------|
| LOG | b_4 |
| LOG | w_1 |
| UNDO | $w_1[b_{56}, \text{cash}=94340.45]$ |
| REDO | $w_1[b_{56}, \text{cash}=84340.45]$ |
| LOG | b_2 |
| UNDO | $w_1[b_{34}, \text{cash}=10900.67]$ |
| REDO | $w_2[b_{34}, \text{cash}=8900.67]$ |
| UNDO | $w_2[b_{67}, \text{cash}=34005.00]$ |
| REDO | $w_2[b_{67}, \text{cash}=36005.25]$ |
| LOG | b_7 |
| LOG | c_7 |
| UNDO | $w_1[b_{34}, \text{cash}=8900.67]$ |
| REDO | $w_1[b_{34}, \text{cash}=18900.67]$ |
| UNDO | $w_7[b_{67}, \text{cash}=36005.25]$ |
| REDO | $w_7[b_{67}, \text{cash}=37005.25]$ |
| LOG | c_7 |
| LOG | c_4 |

Disc Before Recovery

| branch | | |
|-----------------|-------------|----------|
| <u>sortcode</u> | bname | cash |
| 56 | 'Wimbledon' | 84340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

Disc After Recovery

| branch | | |
|-----------------|-------------|----------|
| <u>sortcode</u> | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 37005.25 |

Omitting the REDO Log

Assignment Project Exam Help

If no REDO records kept

must flush committed transactions to data disc

- 1 $C = \emptyset, D = \emptyset$
- 2 Scan the log backwards from the end.
- 3 commit entry \rightarrow add to C
- 4 undo entry for member of $C \rightarrow$ add object to D *without* making changes to the data.
- 5 perform undo entry for object not of member D

Omitting the Undo Log

Assignment Project Exam Help

If no UNDO records kept

transaction must never write uncommitted data

- add **fix** command between RM and CM to stop CM flushing data
- commit is followed by flush or **unfix** of fixed objects

Omitting UNDO and REDO

atomic commit → out of place updating

Quiz 3: Contents of Disc Before Commit if no UNDO log

| branch | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

```

BEGIN TRANSACTION T1
UPDATE branch
SET cash=cash-10000.00
WHERE sortcode=56

UPDATE branch
SET cash=cash+10000.00
WHERE sortcode=34
COMMIT TRANSACTION T1
  
```

| branch ① | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ③ | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 84340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ② | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ④ | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 84340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

What must the contents of the branch table on disc be before the transaction commits?

A

①

B

① or ④

C

④

D

①, ②, ③ or ④

Quiz 4: Contents of Disc After Commit if no REDO log

| branch | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

```
BEGIN TRANSACTION T1
UPDATE branch
SET cash=cash-10000.00
WHERE sortcode=56

UPDATE branch
SET cash=cash+10000.00
WHERE sortcode=34
COMMIT TRANSACTION T1
```

| branch ① | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ③ | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 84340.45 |
| 34 | 'Goodge St' | 8900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ② | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 94340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

| branch ④ | | |
|----------|-------------|----------|
| sortcode | bname | cash |
| 56 | 'Wimbledon' | 84340.45 |
| 34 | 'Goodge St' | 18900.67 |
| 67 | 'Strand' | 34005.00 |

What must the contents of the branch table on disc be after the transaction commits?

A

①

B

① or ④

C


④

D

①, ②, ③ or ④

Checkpointing

Assignment Project Exam Help



<https://powcoder.com>

- Forces the database into some known state
- Recovery limited to only look back to checkpoint (or a 'bit' before!)
 - speeds the recovery operation
 - limits the size of log
- The more consistent this known state
 - the easier it is to recover
 - the longer it takes to perform the checkpoint

Add WeChat powcoder

Commit Consistent Checkpoint

Assignment Project Exam Help

Generating a Commit Consistent Checkpoint

- 1 Stop accepting new transactions
- 2 Finish existing transactions.
- 3 Flush all dirty data & cache objects to disc.
- 4 Write a checkpoint to stable log.

<https://powcoder.com>

Add WeChat powcoder

- recovery now only needs to scan back to *cp* in log ✓
- possible long hold-up at checkpoint ✗

Cache Consistent Checkpoint

Generating a Cache Consistent Checkpoint

- 1 Suspend all transactions
- 2 Flush all dirty cache objects to disc
- 3 Write a checkpoint + active transactions to stable log

Recovery from Cache Consistent Checkpoint records

- 1 perform UNDOs of non-committed transactions back to *cp*
 - 2 perform UNDO of non-committed transactions before *cp* if they were active at *cp*
 - 3 perform REDOs of committed transactions after *cp*
- could still have delay whilst flushing cached objects

Worksheet: Cache Consistent Checkpoint

```

LOG      b7
UNDO     w7[b67, cash=34005.25]
REDO     w7[b67, cash=37005.25]
LOG      b2
UNDO     w2[b34, cash=10900.67]
REDO     w2[b34, cash=8900.67]
LOG      b6
UNDO     w6[a101, rate=5.25]
REDO     w6[a101, rate=6.00]
LOG      b1
UNDO     w1[b56, cash=94340.45]
REDO     w1[b56, cash=84340.45]
LOG      a7
LOG      cp{1, 2, 6}
      ⋮

```

```

      ⋮
UNDO     w6[a119, rate=5.50]
REDO     w6[a119, rate=6.00]
LOG      c6
UNDO     w2[b67, cash=34005.00]
REDO     w2[b67, cash=36005.25]
LOG      b8
LOG      c2
UNDO     w1[b34, cash=8900.67]
REDO     w1[b34, cash=18900.67]
LOG      b9
UNDO     w9[b67, cash=36005.00]
REDO     w9[b67, cash=20000.00]
LOG      c9

```

Fuzzy Checkpointing

Generating a Fuzzy Checkpoint


- 1 Suspend all transactions
- 2 Flush any dirty cache objects to disc not flushed in previous *cp*
- 3 Write a checkpoint + active transactions to stable log

Recovery from Fuzzy Checkpoint records

Recovery works like cache consistent checkpoint, but working with penultimate *cp*

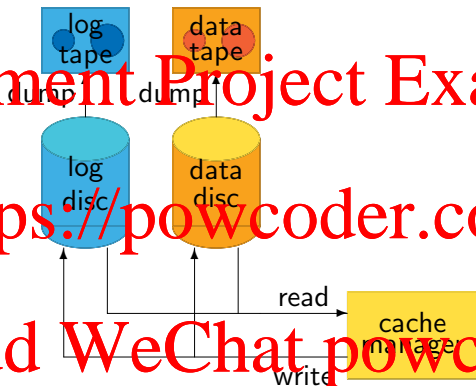
- 1 perform UNDOs of non-committed transactions back to penultimate *cp*
- 2 perform UNDO of non-committed transactions before penultimate *cp* if they were active at *cp*
- 3 perform REDOs of committed transactions after penultimate *cp*

Media Failures: Mirroring (RAID-1)



- Keep more than one active copy of data and log
- Writes sent to both
- Read from either

Media Failures: Dumping



- 'tape' might also be a external file server, removable HD, *etc.*
- To use normal OS backup procedure
 - DBMS must not be still running
 - raw partition must not be used

Checkpoints and Dumps

Assignment Project Exam Help

- Dump must do a checkpoint

- Restore involves:


- 1 copy tape to disc
- 2 undo transactions active at the archive time
- 3 redo transactions that committed after the archive

- commit consistent checkpoint obvious and easy

<https://powcoder.com>

Add WeChat powcoder

Media Failures: Archive Database



- mirror log, but only have one active database
- periodically archive updates onto archive database
- failure of active database disc involves restore of archive database using logs

THE END

Assignment Project Exam Help

- Content of the course is what has been presented in the lectures
- Revise by reviewing worksheets and courseworks
- 2011 exam papers onwards set to current syllabus
- Older exam questions mostly apply, but there is more emphasis on RA and SQL lesson concurrency

<https://powcoder.com>
Add WeChat powcoder